

Outline

- Recap
- Analytical Derivation (Bayes' theorem)
- Variable Elimination
- Belief Propagation

Learning objectives

- Understand the concept of (exact) inference
- Be able to perform analytical derivation inference with Gaussian distribution
- Be able to perform variable elimination with discrete variables
- Be able to perform belief propagation with discrete variables

Previously, in MBML...


- Representation (weeks 1-4)
- Modelling toolbox (weeks 5-9, +13)
- Inference (weeks 10-12)

Previously, in MBML...

- Representation (weeks 1-4)
 - PGM basics
 - Conditional independence
 - Generative processes
 - Joint probability distribution
 - Priors and likelihood
 - Factorization

$$p(\beta, \mathbf{z}, \mathbf{k}, \mathbf{dt}) = p(\beta | \sigma_\beta) \prod_{n=1}^N p(k_n | \mathbf{x}_n, \beta, \sigma_\epsilon) p(z_n | \pi) p(dt_n | z_n, k_n)$$

- 1 Draw a pair of parameters¹, $\beta \sim \mathcal{N}(\mathbf{0}, I\sigma_\beta)$
- 2 For $n = 1..N$
 - 1 Draw one value for z_n , such that $z_n \sim \text{Bern}(\pi)$.
 - If $z_n = 1$, the bus has stopped ($z_n = 0$ otherwise)
 - Distributed as Bernoulli, with parameter π
 - 2 Draw one value for k_n , such that $k_n \sim \mathcal{N}(\mathbf{x}_n^T \beta, \sigma_\epsilon)$
 - 3 If $z_n = 1$, $dt_n = k_n$,
 - otherwise $dt_n = 0$



Previously, in MBML...

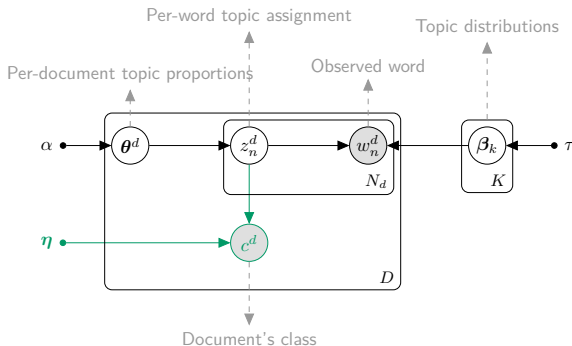
- Representation (weeks 1-4)
- Modelling toolbox (weeks 5-9, +13)
 - Mixture models
 - Different likelihoods (Gaussian, Poisson, etc.)
 - Link functions (log, softmax, Probit, etc.)
 - Non-linear relationships (e.g. using neural networks)
 - Discrete vs. continuous target variables
 - Heteroscedastic models
 - Hierarchical models
 - Temporal models (continuous and discrete)
 - Topic modelling (discrete data; e.g. text corpora)
 - Bayesian non-parametric models (week 13)
- Inference (weeks 10-12)

Previously, in MBML...

- Representation (weeks 1-4)
- Modelling toolbox (weeks 5-9, +13)
 - Bayesian Gaussian Mixture models
 - Bayesian Linear regression
 - Poisson regression
 - Heteroscedastic regression
 - Bayesian Logistic regression
 - Bayesian Probit regression
 - Hierarchical Logistic regression
 - Autoregressive models
 - Linear dynamical systems (e.g. Kalman filter)
 - Hidden Markov models
 - Latent Dirichlet allocation
 - Gaussian processes (week 13)
- Inference (weeks 10-12)

Previously, in MBML...

- Representation (weeks 1-4)
- Modelling toolbox (weeks 5-9, +13)
 - Mix & Match (e.g. LDA + (Logistic) Regression = Supervised LDA)



$$p(c^d | \bar{\mathbf{z}}, \boldsymbol{\eta}) = \frac{\exp(\boldsymbol{\eta}_c^T \bar{\mathbf{z}})}{\sum_{l=1}^C \exp(\boldsymbol{\eta}_l^T \bar{\mathbf{z}})}$$


- Inference (weeks 10-12)

Previously, in MBML...


- Representation (weeks 1-4)
- Modelling toolbox (weeks 5-8, +12)
- Inference (weeks 9-11)
 - Exact inference
 - Analytical (Bayes' rule)
 - Variable elimination
 - Belief propagation
 - Approximate inference
 - Stochastic methods - Markov chain Monte Carlo (MCMC)
 - Deterministic methods - Variational inference (VI)



True distribution



Monte Carlo



Variational

Exact inference

- Computation of the exact posterior probability distribution over the variables of interest
 - Best possible solution, given data and specification
- Analytical derivations:
 - High computational efficiency
 - However, quite often, not possible at all
(when it is not possible, we say that it is *intractable*)
- Algorithmic methods
 - Variable elimination
 - Message passing

Analytical derivation

- The Gaussian form has very nice properties¹
- A multivariate Gaussian with d dimensions is:

$$\mathcal{N}(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{D/2}|\boldsymbol{\Sigma}|^{1/2}} e^{-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1}(\mathbf{x}-\boldsymbol{\mu})}$$

- The inverse of the covariance matrix is $\boldsymbol{\Lambda} = \boldsymbol{\Sigma}^{-1}$ is called *precision matrix*
- If you have a marginal Gaussian distribution for \mathbf{x} and a conditional Gaussian distribution of \mathbf{y} given \mathbf{x} in the form

$$p(\mathbf{x}) = \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Lambda}^{-1})$$

$$p(\mathbf{y}|\mathbf{x}) = \mathcal{N}(\mathbf{y}|\mathbf{Ax} + \mathbf{b}, \mathbf{L}^{-1})$$

- Then the marginal distribution for \mathbf{y} and a conditional Gaussian distribution of \mathbf{x} given \mathbf{y} have the form

$$p(\mathbf{y}) = \mathcal{N}(\mathbf{y}|\mathbf{A}\boldsymbol{\mu} + \mathbf{b}, \mathbf{L}^{-1} + \mathbf{A}\boldsymbol{\Lambda}^{-1}\mathbf{A}^T)$$

$$p(\mathbf{x}|\mathbf{y}) = \mathcal{N}(\mathbf{x}|\boldsymbol{\Gamma}[\mathbf{A}^T\mathbf{L}(\mathbf{y} - \mathbf{b})] + \boldsymbol{\Gamma}\boldsymbol{\mu}, \boldsymbol{\Gamma})$$

- where $\boldsymbol{\Gamma} = (\boldsymbol{\Lambda} + \mathbf{A}^T\mathbf{L}\mathbf{A})^{-1}$

¹ Check appendix of Bishop's book for many more useful properties!

Analytical derivation

- If we have a joint Gaussian distribution $N(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma})$, with $\boldsymbol{\Lambda} = \boldsymbol{\Sigma}^{-1}$, and we define the following partitions:

$$\mathbf{x} = \begin{pmatrix} \mathbf{x}_a \\ \mathbf{x}_b \end{pmatrix}, \boldsymbol{\mu} = \begin{pmatrix} \boldsymbol{\mu}_a \\ \boldsymbol{\mu}_b \end{pmatrix}$$

$$\boldsymbol{\Sigma} = \begin{pmatrix} \boldsymbol{\Sigma}_{aa} & \boldsymbol{\Sigma}_{ab} \\ \boldsymbol{\Sigma}_{ba} & \boldsymbol{\Sigma}_{bb} \end{pmatrix}, \boldsymbol{\Lambda} = \begin{pmatrix} \boldsymbol{\Lambda}_{aa} & \boldsymbol{\Lambda}_{ab} \\ \boldsymbol{\Lambda}_{ba} & \boldsymbol{\Lambda}_{bb} \end{pmatrix}$$

- Then the conditional distribution $p(\mathbf{x}_a|\mathbf{x}_b)$ is given by

$$p(\mathbf{x}_a|\mathbf{x}_b) = \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_{a|b}, \boldsymbol{\Lambda}_{aa}^{-1})$$

$$\boldsymbol{\mu}_{a|b} = \boldsymbol{\mu}_a - \boldsymbol{\Lambda}_{aa}^{-1} \boldsymbol{\Lambda}_{ab}(\mathbf{x}_b - \boldsymbol{\mu}_b)$$

- And the marginal distribution $p(\mathbf{x}_a)$ is given by

$$p(\mathbf{x}_a) = \mathcal{N}(\mathbf{x}_a|\boldsymbol{\mu}_a, \boldsymbol{\Sigma}_{aa})$$

Analytical derivation

- If we have two (univariate) Gaussian distributions (for the **same** variable x):

$$p(x) = \mathcal{N}(x|\mu_a, \sigma_a^2), \quad p(x) = \mathcal{N}(x|\mu_b, \sigma_b^2)$$

- Their product is:

$$p(x) = \mathcal{N}(x|\mu_{ab}, \sigma_{ab}^2)$$

- where

$$\mu_{ab} = \frac{\mu_a \sigma_b^2 + \mu_b \sigma_a^2}{\sigma_a^2 + \sigma_b^2}, \quad \sigma_{ab}^2 = \frac{\sigma_a^2 \sigma_b^2}{\sigma_a^2 + \sigma_b^2}$$

- This is very useful to directly combine models into a single prediction (e.g. ensemble models)!

Analytical derivation

- Example of Bayesian linear regression
- The joint probability of our model is given by:

$$p(\mathbf{y}, \boldsymbol{\beta} | \mathbf{X}, \lambda, \sigma) = p(\boldsymbol{\beta} | \mathbf{0}, \lambda \mathbf{I}) \prod_{n=1}^N p(y_n | \boldsymbol{\beta}^T \mathbf{x}_n, \sigma^2)$$

- Since both our prior $p(\boldsymbol{\beta} | \lambda)$ and likelihood $p(y_n | \mathbf{x}_n, \boldsymbol{\beta}, \sigma)$ are Gaussian, we apply Bayes' theorem to compute posterior over $\boldsymbol{\beta}$:

$$\begin{aligned} \underbrace{p(\boldsymbol{\beta} | \mathbf{y}, \mathbf{X}, \lambda, \sigma)}_{\text{posterior}} &\propto \underbrace{p(\boldsymbol{\beta} | \mathbf{0}, \lambda \mathbf{I})}_{\text{prior}} \underbrace{\prod_{n=1}^N p(y_n | \boldsymbol{\beta}^T \mathbf{x}_n, \sigma^2)}_{\text{likelihood}} \\ &= \mathcal{N}(\boldsymbol{\beta} | \mathbf{0}, \lambda \mathbf{I}) \prod_{n=1}^N \mathcal{N}(y_n | \boldsymbol{\beta}^T \mathbf{x}_n, \sigma^2) \\ &= \mathcal{N}(\boldsymbol{\beta} | \boldsymbol{\mu} = ?, \boldsymbol{\Sigma} = ?) \end{aligned}$$

- Can you do it? $\boldsymbol{\mu} = ?$, $\boldsymbol{\Sigma} = ?$ (Hint: look at slide 11...)

Analytical derivation

$$\begin{aligned}
 \underbrace{p(\boldsymbol{\beta}|\mathbf{y}, \mathbf{X}, \lambda, \sigma)}_{\text{posterior}} &\propto \underbrace{p(\boldsymbol{\beta}|\mathbf{0}, \lambda\mathbf{I})}_{\text{prior}} \underbrace{\prod_{n=1}^N p(y_n|\boldsymbol{\beta}^T \mathbf{x}_n, \sigma^2)}_{\text{likelihood}} \\
 &= \mathcal{N}(\boldsymbol{\beta}|\mathbf{0}, \lambda\mathbf{I}) \prod_{n=1}^N \mathcal{N}(y_n|\boldsymbol{\beta}^T \mathbf{x}_n, \sigma^2)
 \end{aligned}$$

- Notice that, because our observations y_n are i.i.d., we can re-write:

$$\begin{aligned}
 &= \mathcal{N}(\boldsymbol{\beta}|\mathbf{0}, \lambda\mathbf{I}) \mathcal{N}(\mathbf{y}|\mathbf{X}\boldsymbol{\beta}, \sigma^2\mathbf{I}) \\
 &= \mathcal{N}(\boldsymbol{\beta}|\boldsymbol{\mu}=?, \boldsymbol{\Sigma}=?)
 \end{aligned}$$

- where $\mathbf{y} = \{y_1, \dots, y_N\}$ and $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$
- Thus, we can apply the properties in slide 11, yielding:

$$\begin{aligned}
 \boldsymbol{\mu} &= \boldsymbol{\Sigma} (\sigma^{-2} \mathbf{X}^T \mathbf{y}) \\
 \boldsymbol{\Sigma} &= (\lambda^{-1} \mathbf{I} + \sigma^{-2} \mathbf{X}^T \mathbf{X})^{-1}
 \end{aligned}$$

Variable Elimination

simple PGM

- A chain graph



$$p(z_1, z_2, z_3, z_4) = p(z_4|z_3) p(z_3|z_2) p(z_2|z_1) p(z_1)$$

- Notice that: $z_4 \perp\!\!\!\perp \{z_1, z_2\} | z_3$
- We want to make inference on z_4

$$p(z_4) = \sum_{z_1, z_2, z_3} p(z_1, z_2, z_3, z_4)$$

Variable Elimination

simple PGM



- A trivial solution is to go over all possible combinations of values!

$$p(z_4) = \sum_{z_1} \sum_{z_2} \sum_{z_3} p(z_1, z_2, z_3, z_4)$$

- Generally, if each of the m variables has k possible values, we'd have a complexity of $O(k^{m-1})$
- This quickly becomes intractable (just remember our *trivial* example with a mixture model) \rightarrow NP-hard problem

Variable Elimination

simple PGM

- We can take advantage of the PGM structure

$$\begin{aligned} p(z_4) &= \sum_{z_1, z_2, z_3} p(z_1, z_2, z_3, z_4) \\ &= \sum_{z_1, z_2, z_3} p(z_4|z_3) p(z_3|z_2) p(z_2|z_1) p(z_1) \\ &= \sum_{z_3} p(z_4|z_3) \sum_{z_2} p(z_3|z_2) \sum_{z_1} p(z_2|z_1) p(z_1) \end{aligned}$$

- In a chain graph, the complexity reduces to $O(mk^2)$

Variable Elimination

simple PGM



$$p(z_4) = \sum_{z_3} p(z_4|z_3) \sum_{z_2} p(z_3|z_2) \sum_{z_1} p(z_2|z_1) p(z_1)$$

Variable Elimination

simple PGM




$$\begin{aligned} p(z_4) &= \sum_{z_3} p(z_4|z_3) \sum_{z_2} p(z_3|z_2) \sum_{z_1} p(z_2|z_1) p(z_1) \\ &= \sum_{z_3} p(z_4|z_3) \sum_{z_2} p(z_3|z_2) f_a(z_2) \end{aligned}$$

- Notice that $f_a(z_2)$ is a function of z_2 (also called a *factor*). For example, a CPT with the probabilities of the k values of z_2
- We just “got rid” of z_1 by marginalizing over its values
- Same as in last lecture, when we marginalized over z for implementing LDA in STAN...

Variable Elimination

simple PGM



$$\begin{aligned} p(z_4) &= \sum_{z_3} p(z_4|z_3) \sum_{z_2} p(z_3|z_2) \sum_{z_1} p(z_2|z_1) p(z_1) \\ &= \sum_{z_3} p(z_4|z_3) \sum_{z_2} p(z_3|z_2) f_a(z_2) \\ &= \sum_{z_3} p(z_4|z_3) f_b(z_3) \end{aligned}$$

Variable Elimination

simple PGM



$$\begin{aligned} p(z_4) &= \sum_{z_3} p(z_4|z_3) \sum_{z_2} p(z_3|z_2) \sum_{z_1} p(z_2|z_1) p(z_1) \\ &= \sum_{z_3} p(z_4|z_3) \sum_{z_2} p(z_3|z_2) f_a(z_2) \\ &= \sum_{z_3} p(z_4|z_3) f_b(z_3) \\ &= f_c(z_4) \end{aligned}$$

Variable Elimination (VE)


- Time complexity is exponential in size of largest factor
 - Each f_i is a factor
 - Size of factor is number of variables it depends on
- Order is vital (and sometimes a complicated problem)!
- Observed variables

$$p(z|x = k) = \frac{p(z, x = k)}{p(x = k)}$$

- Perform VE on $p(z, x = k)$ (i.e. we fix $x=k$) and then on $p(x = k) = \sum_z p(z, x = k)$
- **Only acyclic graphs!**

Playtime!

- Apply the Variable Elimination algorithm to the following graph
- We want to infer $p(d|b = 1)$



	$p(a b,c)$	
	$a = 0$	$a = 1$
$b = 0, c = 0$	0.7	0.3
$b = 0, c = 1$	0.3	0.7
$b = 1, c = 0$	0.5	0.5
$b = 1, c = 1$	0.1	0.9

$p(c)$	
$c = 0$	$c = 1$
0.7	0.3

$p(b)$	
$b = 0$	$b = 1$
0.4	0.6

$p(d a)$		
	$d = 0$	$d = 1$
$a = 0$	0.6	0.4
$a = 1$	0.2	0.8

Polytrees

- A graph is a polytree if (and only if) there is at most one simple path between any two nodes, v_i and v_k




Figure: Polytree




Figure: ?




Figure: Not polytree





Figure: ?

Belief propagation


Pearl's algorithm

- We want to solve $p(\mathbf{z}|\mathbf{x})$, i.e. the conditional probability of all variables $\mathbf{z} = \{z_1, z_2, \dots, z_V\}$, given evidence $\mathbf{x} = \{x_1, x_2, \dots, x_E\}$
- Graph structure allows for incremental inference steps
 - Like in Variable Elimination...
- Some nodes are clearly specified from beginning
 - Evidence $x_i = k$
 - Priors
- The other nodes can build on them
- Information flows as **messages** between nodes


Belief propagation intuition




Belief propagation intuition




Belief propagation intuition




Belief propagation intuition



Belief propagation intuition




Belief propagation intuition



Belief propagation


Pearl's algorithm



- π messages - from parent to child

Belief propagation

Pearl's algorithm




- π messages - from parent to children
- Represent the current belief about the parent - **causal evidence**
- After receiving all parents' messages, one can know more of the child

$$\pi(z_j) = \sum_{z_i \in \text{parent}(z_j)} p(z_j | z_1, z_2, \dots) \prod_{z_i \in \text{parent}(z_j)} \pi_{z_i, z_j}(z_i)$$

Belief propagation

Pearl's algorithm




- π **messages** - from parent to children
- Represent the current belief about the parent - **causal evidence**
- After receiving all parents' messages, one can know more of the child

factorization: $p(z_j, z_1, z_2, \dots, z_j) =$
 $p(z_j | z_1, z_2, \dots, z_j) p(z_1), p(z_2), \dots, p(z_j)$

$$p(z_j) = \sum_{z_1, z_2, \dots, z_j} p(z_j | z_1, z_2, \dots, z_j) p(z_1), p(z_2), \dots, p(z_j)$$

Belief propagation

Pearl's algorithm




- π **messages** - from parent to children
- Represent the current belief about the parent - **causal evidence**
- After receiving all parents' messages, one can know more of the child

$$\pi(z_j) = \sum_{z_i \in \text{parent}(z_j)} p(z_j | z_1, z_2, \dots) \prod_{z_i \in \text{parent}(z_j)} \pi_{z_i, z_j}(z_i)$$

Belief propagation

Pearl's algorithm



- **λ messages** - from child to parent

$$p(z_j | z_i)$$

- Represents the belief about the parent's value - **diagnostic evidence**


Belief propagation

Pearl's algorithm

- **λ messages** - from child to parent

$$p(z_i | z_j)$$

- Represents the belief about the parent's value - **diagnostic evidence**
- After receiving all children's messages, one can know more of the parent




$$\lambda(z_j) = \prod_{z_i \in \text{child}(z_j)} \lambda_{z_i, z_j}(z_j)$$

Belief propagation

Pearl's algorithm


- But we still need the formulas for the messages!

$$\pi_{z_i, z_j}(z_i) = \pi(z_i) \prod_{v \neq j} \lambda_{z_v, z_i}(z_i)$$



- λ message when we have $p(z_i|z_j)$ (only one parent, z_j)

$$\lambda_{z_i, z_j}(z_j) = \sum_{z_i} \lambda(z_i) p(z_i|z_j)$$




Belief propagation

Pearl's algorithm

- Generic formula for λ messages:

$$\lambda_{z_i, z_j}(z_j) = \sum_{z_i} \lambda(z_i) \sum_{z_v \in \mathbf{z} \setminus z_j} p(z_i | z_1, z_2, \dots, z_j) \prod_{v \neq j} \pi_{z_v, z_i}(z_v)$$



Pearl's algorithm - Initialization step


- For all $x_i \in \mathbf{x}$ (evidence nodes):
 - $\lambda(x_i) = 1$, wherever $x_i = e_i$, 0 otherwise (i.e. change the CPT)
 - $\pi(x_i) = 1$, wherever $x_i = e_i$, 0 otherwise (i.e. change the CPT)
- For all nodes, z_i , without parents:
 - $\pi(z_i) = p(z_i)$, the prior
- For all nodes, z_i , without children:
 - $\lambda(z_i) = 1$

Pearl's algorithm

- Iterate until no change occurs
 - 1 For each node z_i , if it has received all π_{*,z_i} messages, calculate $\pi(z_i)$
 - 2 For each node z_i , if it has received all λ_{*,z_i} messages, calculate $\lambda(z_i)$
 - 3 For each node z_i , if $\pi(z_i)$ is calculated, and all λ_{*,z_i} messages have been received (except from z_j), calculate $\pi_{z_i,z_j}(z_i)$ and send the message to z_j
 - 4 For each node z_i , if $\lambda(z_i)$ is calculated, and all π_{*,z_i} messages have been received (except from z_j), calculate $\lambda_{z_i,z_j}(z_j)$ and send the message to z_j
- Compute $\text{Bel}(z_i) \propto \lambda(z_i) \pi(z_i)$ and normalize, for all desired nodes

Example

- Let's apply Belief Propagation to the following graph
- We want to infer $p(d|b = 1)$



	$p(a b,c)$	
	$a = 0$	$a = 1$
$b = 0, c = 0$	0.7	0.3
$b = 0, c = 1$	0.3	0.7
$b = 1, c = 0$	0.5	0.5
$b = 1, c = 1$	0.1	0.9

$p(c)$	
$c = 0$	$c = 1$
0.7	0.3

$p(b)$	
$b = 0$	$b = 1$
0.4	0.6

$p(d a)$		
	$d = 0$	$d = 1$
$a = 0$	0.6	0.4
$a = 1$	0.2	0.8

Example

- From the initialization, we have

$\pi(c) = p(c)$	
$c = 0$	$c = 1$
0.7	0.3


$\pi(b)$	
$b = 0$	$b = 1$
0	1

$\lambda(b)$	
$b = 0$	$b = 1$
0	1

$\lambda(d)$	
$b = 0$	$b = 1$
1	1

Example

- We start with the evidence and prior



$$\pi_{c,a}(c) = \pi(c) \prod_{x \neq a} \lambda_{x,c}(c) = \pi(c)$$


$\pi_{c,a}(c)$	
$c = 0$	$c = 1$
0.7	0.3

$$\pi_{b,a}(b) = \pi(b)$$

$\pi_{b,a}(b)$	
$b = 0$	$b = 1$
0	1

Example

- Now, $\pi(a)$




$$\pi(a) = \sum_{b,c} p(a|b,c) \pi_{b,a}(b) \pi_{c,a}(c)$$

For $a=1$:

$$\begin{aligned}
 \pi(a=1) &= \sum_{b=\{0,1\}} \sum_{c=\{0,1\}} p(a=1|b,c) \pi_{b,a}(b) \pi_{c,a}(c) \\
 &= \sum_{c=\{0,1\}} p(a=1|b=1,c) \pi_{c,a}(c) \\
 &= p(a=1|b=1,c=0) \pi_{c,a}(c=0) + p(a=1|b=1,c=1) \pi_{c,a}(c=1) \\
 &= 0.5 \times 0.7 + 0.9 \times 0.3 = 0.62, \text{ so } \pi(a=1) = 0.62 \text{ and } \pi(a=0) = 0.38
 \end{aligned}$$

Example

- Finally, we reach d



Since a has no other children, we have $\pi_{a,d}(a) = \pi(a)$

$$\pi_{a,d}(a)$$

$\pi_{a,d}(a)$	
$a = 0$	$a = 1$
0.38	0.62

Example

$$\pi(d) = \sum_{a=\{0,1\}} p(d|a) \pi_{a,d}(a)$$

$p(d a)$		
	$d = 0$	$d = 1$
$a = 0$	0.6	0.4
$a = 1$	0.2	0.8

$\pi_{a,d}(a)$		
	$a = 0$	$a = 1$
	0.38	0.62

- Trivially becomes:


$\pi(d)$		
	$b = 0$	$b = 1$
	0.352	0.648

- $\text{Bel}(d) = \pi(d)\lambda(d)$
- In this case, the solution is trivially:

$$p(d|b=1) = \pi(d)$$

- But why stop here? We can propagate back, and also get $p(a|b=1)$ and $p(c|b=1)$

Example



$$\lambda_{d,a}(a) = \sum_{d=\{0,1\}} \lambda(d)p(d|a)$$


- Since $\lambda(d) = 1$, we have:

$$\lambda_{d,a}(a) = \sum_{d=\{0,1\}} p(d|a) = 1$$

- Notice that we have a unnormalized table!

$\lambda_{d,a}(a)$	
$a = 0$	$a = 1$
1	1

Example




- Since $\lambda(a) = 1$, we have:

$$\lambda(c) = \lambda_{a,c}(c) = \sum_{a=\{0,1\}} \lambda(a) p(a|b=1, c) = \sum_{a=\{0,1\}} p(a|b=1, c) = 1$$

- Therefore:
 - $\text{Bel}(c) = \pi(c) \lambda(c) = \pi(c) = p(c)$
 - $\text{Bel}(a) = \pi(a) \lambda(a) = \pi(a) = \sum_c p(a|b=1, c) \pi_{c,a}(c)$, as calculated before
- Makes sense, right (notice the independence of the graph!)?

Example

- But, what if there's a new evidence, $d = 0$?




- We only need to update d and propagate back!
- A new initialization gives:

$\lambda(d) = \pi(d)$	
$d = 0$	$d = 1$
1	0

$$\lambda_{d,a}(a) = \sum_{d=\{0,1\}} \lambda(d)p(d|a) = p(d=0|a)$$

Example

- But, what if there's a new evidence, $d = 0$?



- Tables we need:

	$p(d a)$	
	$d = 0$	$d = 1$
$a = 0$	0.6	0.4
$a = 1$	0.2	0.8

$\lambda(a)$	
$a = 0$	$a = 1$
0.6	0.2


$$\lambda_{d,a}(a) = \sum_{d=\{0,1\}} \lambda(d)p(d|a) = p(d=0|a)$$

$$\lambda(a) = \lambda_{d,a}(a)$$

$$\begin{aligned} \lambda_{a,c}(c) &= \sum_{a,b=\{0,1\}} \lambda(a)p(a|b,c)\pi_{b,a}(b) = \\ &= \sum_{a=\{0,1\}} \lambda(a)p(a|b=1,c) \end{aligned}$$

Example

- But, what if there's a new evidence, $d = 0$?



- Tables we need:

	$p(a b,c)$	
	$a = 0$	$a = 1$
$b = 0, c = 0$	0.7	0.3
$b = 0, c = 1$	0.3	0.7
$b = 1, c = 0$	0.5	0.5
$b = 1, c = 1$	0.1	0.9


$\lambda(a)$	
$a = 0$	$a = 1$
0.6	0.2

$$\lambda_{d,a}(a) = \sum_{d=\{0,1\}} \lambda(d)p(d|a) = p(d=0|a)$$

$$\lambda(a) = \lambda_{d,a}(a)$$

$$\begin{aligned}
 \lambda_{a,c}(c) &= \sum_{a,b=\{0,1\}} \lambda(a)p(a|b,c)\pi_{b,a}(b) = \\
 &= \sum_{a=\{0,1\}} \lambda(a)p(a|b=1,c) \\
 &= 0.6 \times p(a=0|b=1,c) + 0.2 \times p(a=1|b=1,c)
 \end{aligned}$$

Example



- Tables we need:

	$p(a b,c)$	
	$a = 0$	$a = 1$
$b = 0, c = 0$	0.7	0.3
$b = 0, c = 1$	0.3	0.7
$b = 1, c = 0$	0.5	0.5
$b = 1, c = 1$	0.1	0.9

$$\begin{aligned}
 \lambda_{a,c}(c) &= \sum_{a,b=\{0,1\}} \lambda(a)p(a|b,c)\pi_{b,a}(b) = \\
 &= \sum_{a=\{0,1\}} \lambda(a)p(a|b=1,c) \\
 &= 0.6 \times p(a=0|b=1,c) + 0.2 \times p(a=1|b=1,c)
 \end{aligned}$$

For $c = 0$

$$\begin{aligned}
 \lambda_{a,c}(c = 0) &= 0.6 \times p(a=0|b=1,c=0) + 0.2 \times p(a=1|b=1,c=0) \\
 &= 0.6 \times 0.5 + 0.2 \times 0.5 = 0.4
 \end{aligned}$$

For $c = 1$

$$\lambda_{a,c}(c = 1) = 0.6 \times 0.1 + 0.2 \times 0.9 = 0.24$$

$\lambda(a)$	
$a = 0$	$a = 1$
0.6	0.2

Example

- So,
- $\text{Bel}(a) = \alpha \lambda(a) \pi(a)$, with α the normalizing factor

$\pi(a)$	
$a = 0$	$a = 1$
0.38	0.62

$\lambda(a)$	
$a = 0$	$a = 1$
0.6	0.2

- $\alpha = (0.38 * 0.6 + 0.62 * 0.2)^{-1} = 2.84$

$$p(a|b = 1, d = 0) = \text{Bel}(a)$$

$a = 0$	$a = 1$
0.65	0.35


$$p(c|b = 1, d = 0) = \text{Bel}(c)$$

$c = 0$	$c = 1$
0.8	0.2

Belief Propagation

- A *forward-backward* pass of the BP algorithm gives us the exact inference for every variable
- Updating can be done incrementally
- Plates can be treated by expansion (given that it is still a polytree)
- We used discrete variables, but the reasoning is the same for continuous
 - Summations become integrals
 - Need to derive a new function for each step (easy in some cases, particularly with exponential family)
- All of the above is valid for polytrees

Playtime!



$p(c)$	
$c = 0$	$c = 1$
0.7	0.3

$p(b a)$		
	$b = 0$	$b = 1$
$a = 0$	0.3	0.7
$a = 1$	0	1



$p(a c,d)$		
	$a = 0$	$a = 1$
$c = 0, d = 0$	0.5	0.5
$c = 0, d = 1$	0.9	0.1
$c = 1, d = 0$	0.1	0.9
$c = 1, d = 1$	0	1

$p(d)$	
$d = 0$	$d = 1$
0.8	0.2

$p(e a)$		
	$e = 0$	$e = 1$
$a = 0$	0.2	0.8
$a = 1$	0.7	0.3



- Consider the graph and CPTs above
- Perform forward-backward Belief Propagation, assuming the evidence that $d = 1, e = 1$

The problem of non-polytrees



- Creates an (infinite) cycle
 - Loopy belief propagation
 - Clique Trees

The problem of non-polytrees



- Creates an (infinite) cycle
 - **Loopy belief propagation**
 - Clique Trees

Loopy Belief Propagation

- Apply BP on the original graph in the following way:
 - ① Initialize all messages to 1
 - ② Run BP algorithm as before (start with evidence/priors)
 - ③ Each node sends messages in parallel (i.e. when it sends to one child/parent, it sends to all children/parents)
 - ④ Loop until convergence
- Works very well when there are some loops, but not a fully connected graph
- With one complete loop, the MAP should be correct
- If $p(\mathbf{z})$ is jointly Gaussian, Loopy Belief Propagation will converge to the correct marginals

Some final notes

- There are actually many tools that do BP
 - Just check: <https://www.cs.ubc.ca/~murphyk/Software/bnsoft.html>
- BP is often applied on the tractable sub-parts of your model, for example combining belief from different sub-models:
 - A Random Forest, RF provides $p_{\text{RF}}(y)$, and a Logistic Regression model, LR, provides $p_{\text{LR}}(y)$:

$$\text{Bel}(y) = \alpha p_{\text{RF}}(y) p_{\text{LR}}(y)$$

- Also called *Bayesian Model Averaging*
- If you have gentle analytical forms for the sub-models (e.g. normal distribution for outputs), the combination is straightforward
- You can even treat sub-models as “priors” that provide to your random variables, and combine many in a complex model!

References

- Kautz and Lowd, CSE 573: Introduction to Artificial Intelligence.
https://courses.cs.washington.edu/courses/cse573/05au/10_26_pearl.ppt
- Koller and Friedman, (2009) Koller, D., and Friedman, N. Probabilistic graphical models: principles and techniques. MIT press. (2009) (Chapters 9 and 10).
- Christopher, M. Bishop. Pattern Recognition and Machine Learning. Springer-Verlag New York, 2016.